

CHETAN PRAKASH

chetan78564@gmail.com — +91 7838083553 — LinkedIn: Chetan Prakash — GitHub: Chetan4124

Summary

Final-year B.Tech (AI & ML) student specializing in Generative AI, LLMs, and full-stack AI development. Experienced in building production-ready RAG pipelines, LangChain agents, and transformer-based models integrated with Django REST Framework and React.js. Passionate about solving real-world problems through scalable AI systems.

Education

J.C. Bose University of Science and Technology, YMCA	Faridabad, India
B.Tech – Computer Science and Engineering (AI & ML) — CGPA: 7.8/10	2022 – 2026
Cambridge Senior Secondary School (CBSE)	Dehradun, India
SSSC: 85%	2020 – 2022

Skills

Languages: Python, JavaScript, C, C++

Libraries & Frameworks: Scikit-learn, Pandas, NumPy, Matplotlib, Seaborn, LangChain, FastAPI, PyTorch, Node.js, React.js, Express.js, Django REST Framework

AI & ML: Machine Learning, Deep Learning, Generative AI, Prompt Engineering, LLMs, RAG, Fine-Tuning, Transfer Learning, Embeddings, Vector Databases, Model Evaluation

Tools & Platforms: Jupyter, VS Code, OpenAI API, Pinecone, FAISS, Chroma, MySQL, Power BI, Tableau, Docker, Git, CI/CD

Methodologies: Agile, REST API Design, Microservices

Work Experience

AI/ML & Django Intern – NIELIT, Delhi *Ongoing*

- Deployed AI/ML models with DRF backends, cutting average API response time by ~30% via optimized inference pipelines.
- Built 3+ scalable REST APIs serving LLM-powered features using Python, deep learning, and LangChain frameworks.
- Managed end-to-end Agile workflows from data preprocessing to production deployment via CI/CD pipelines, reducing iteration cycles by ~25%.

Machine Learning Intern – Suvidha Foundation

Oct 2025 – Nov 2025

- Automated 5+ data workflows using Python & Pandas, reducing manual processing time by ~40%.
- Integrated OpenAI GPT-3 API for text generation and summarization, improving content turnaround by ~3×.
- Delivered a transformer-based text summarizer achieving ~85% ROUGE-1 score, deployed as a standalone module.

Projects

Echo_124M: A 124M Parameter LLM *Python, PyTorch, Hugging Face, NLP, Transformer Architecture*

- Architected a GPT-style decoder-only LLM (~124M params) from scratch, achieving stable convergence within 10 epochs on large text corpora.
- Implemented multi-head self-attention, positional embeddings, and BPE tokenization, reducing overhead by ~20% over baseline.
- Achieved perplexity ~45 on validation data via LR scheduling and dropout regularization; deployed on Hugging Face Spaces for low-latency inference with support for fine-tuning.

Legal_Saathi: AI-Based Legal Advisor

Python, LangChain, DRF, React.js, RAG, LLM Agents

- Built a LangChain RAG-powered legal advisory platform using FAISS vector database, achieving ~90% query relevance accuracy.
- Designed a DRF REST API handling 10+ concurrent sessions with JWT auth and sub-200ms response time.
- Developed a React.js frontend reducing query-to-answer time by ~60% vs. manual legal research.

AI-Based Quiz Generator

Python, LangChain, DRF, React.js, RAG, LLM Agents

- Built an AI quiz platform using LangChain agents to generate 10–20 topic-specific questions/session with ~88% relevance.
- DRF backend supports full quiz lifecycle (create, attempt, evaluate) serving 5+ concurrent users at <300ms latency.

Certificates

Machine Learning Specialization – IBM	Sep 2025	Deep Learning with Keras & TensorFlow	Sep 2025
Web Development Specialization – IBM	Jul 2025	Deep Learning with PyTorch – IBM	Sep 2025
Data Science Specialization – IBM	Jul 2025	IBM AI Developer	Jul 2025
Prompt Specialization – Google	Jul 2025	React Specialization	Jul 2025
Data Visualization – Microsoft	Jul 2025		